

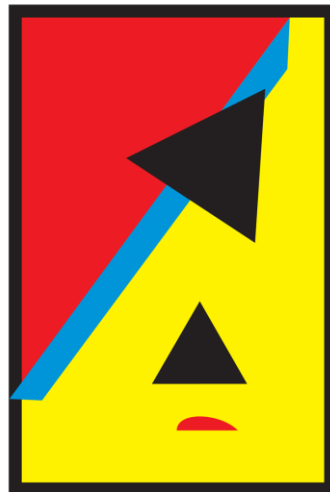
C++ Software Engineering

for engineers of other disciplines

Module 10

"C++ Parallelism"

4th Lecture: Multiprocessing



ALTE N

Spring 2022

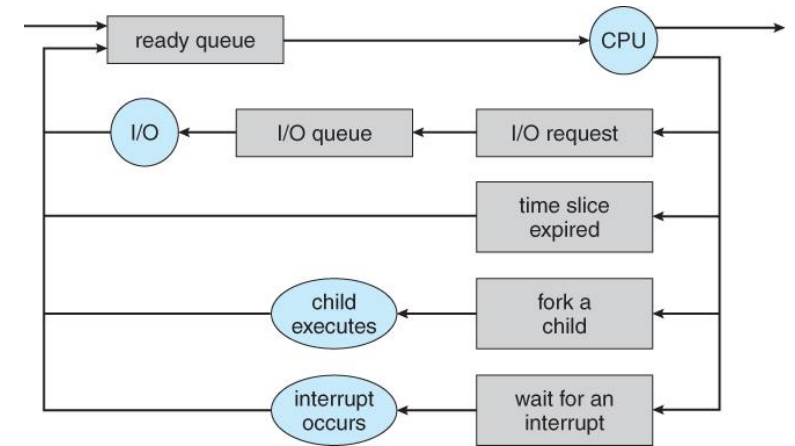

Gothenburg, Sweden

Process

- Process is a running executable which consumes resources *through* operating system interfaces – both APIs and ABIs.
- Unix systems inherit **sysV** APIs.
- **POSIX** is an effort by IEEE to standardized OS interfaces for the purpose of portability.
- Processes can execute instructions, request I/O, wait for an interrupt or spawn a child.

In computer software, an application binary interface (ABI) is an interface between two binary program modules; often, one of these modules is a library or operating system facility, and the other is a program that is being run by a user.

https://en.wikipedia.org/wiki/Application_binary_interface



https://www.cs.uic.edu/~jbelle/CourseNotes/OperatingSystems/3_Processes.html

Fork

- Child process is a copy of its parent except:
 - It has its own unique *process identification* a.k.a. **pid**.
 - Its parents **pid** is different – obviously!
 - It is a fresh copy! (empty queues, *reseted* flags for *locks*, timers, fcntl, ...)
- **fork** is the function to create a child process.
- Children should report their return status to the parent, or they become a zombie process – it is the parent responsibility to **wait** for its children.
- Defunct process are those which have been terminated but their *exit status* has not been read by their parents – POSIX provides *automatic reaping* mechanism.

```
signal(SIGCHLD, SIG_IGN);
```

```
pid_t pid;  
for (size_t i=0 ; i < 2 ; i++)  
    pid = fork();
```

FORK(2)

Linux Programmer's Manual

NAME [top](#)

fork - create a child process

SYNOPSIS [top](#)

```
#include <sys/types.h>  
#include <unistd.h>
```

```
pid_t fork(void);
```


"The entire virtual address space of the parent is replicated in the child, including the states of mutexes, condition variables, and other pthreads objects; the use of **pthread_atfork** may be helpful for dealing with problems that this can cause."

<https://man7.org/linux/man-pages/man2/fork.2.html>

Fork



© M. Rashid Zamani




```
pid_t pid;  
for (size_t i=0 ; i < 2 ; i++)  
    pid = fork();
```

P 0

- Child processes spawned using **fork** returns zero on the child, the returns the child's pid on the parent call.

Fork

```
pid_t pid;  
for (size_t i=0 ; i < 2 ; i++)  
    pid = fork();
```



- A *fresh* copy of the process with the exact content of the virtual memory is *prepared* for the child.


Fork

→

```
pid_t pid;  
for (size_t i=1 ; i < 2 ; i++)  
    pid = fork();
```

→

```
pid_t pid;  
for (size_t i=1 ; i < 2 ; i++)  
    pid = fork();
```




Fork

→

```
pid_t pid;  
for (size_t i=1 ; i < 2 ; i++)  
    pid = fork();
```

→

```
pid_t pid;  
for (size_t i=1 ; i < 2 ; i++)  
    pid = fork();
```



Fork

→

```
pid_t pid;  
for (size_t i=2 ; i < 2 ; i++)  
    pid = fork();
```

→


```
pid_t pid;  
for (size_t i=2 ; i < 2 ; i++)  
    pid = fork();
```

→

```
pid_t pid;  
for (size_t i=2 ; i < 2 ; i++)  
    pid = fork();
```

→

```
pid_t pid;  
for (size_t i=2 ; i < 2 ; i++)  
    pid = fork();
```

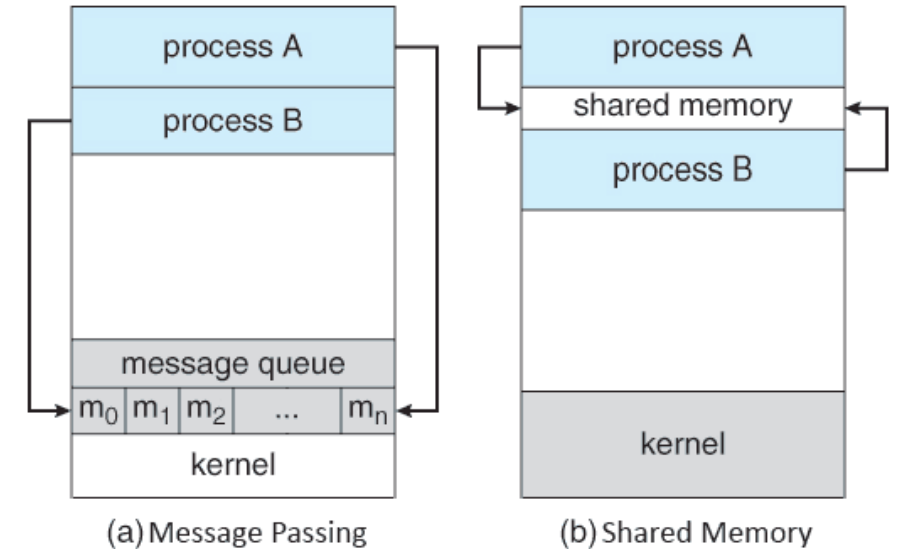


Process Tree -- 2^n

- Child processes spawned using **fork** will form a tree – children's return exit shall be read by their parents.

Inter Process Communication

- IPC is the mechanism used to communicate between processes.
- Process can *signal* each other or *share* data with each other.
- This communication is in Operating System *realm*.
- There are libraries and framework who have implemented IPC.



<https://networkencyclopedia.com/wp-content/uploads/2019/09/inter-process-communications-models.png>

```
#include <string>
#include <zmq.hpp>
int main()
{
    zmq::context_t ctx;
    zmq::socket_t sock(ctx, zmq::socket_type::push);
    sock.bind("inproc://test");
    const std::string_view m = "Hello, world";
    sock.send(zmq::buffer(m), zmq::send_flags::dontwait);
}
```

<https://zeromq.org/languages/cplusplus/>



Chapter 18. Boost.Interprocess

https://www.boost.org/doc/libs/1_74_0/doc/html/interprocess.html

- Anything involving kernel is slow: message queues, signals, and etc. Pipes usually are used to redirect input outputs to each other, FIFO (or file pipes?) are named pipes.
- Resources used between processes needs synchronization.
- Many IPCs happen through *fds* – they need to lock!
- The best is to share memory between processes:
 - Some could argue Sockets are good as well!
 - A shared memory between processes is almost identical to the process's memory – zero overhead!
- Both **sysV** and **POSIX** provide API for that – semaphore is used for resource synchronization over shared memory.

```
struct flock
{
    short int l_type; /* Type of lock: F_RDLCK, F_WRLCK, or F_UNLCK. */
    short int l_whence; /* Where 'l_start' is relative to (like 'lseek'). */
#ifdef __USE_FILE_OFFSET64
    __off_t l_start; /* Offset where the lock begins. */
    __off_t l_len; /* Size of the locked area; zero means until EOF. */
#else
    __off64_t l_start; /* Offset where the lock begins. */
    __off64_t l_len; /* Size of the locked area; zero means until EOF. */
#endif
    __pid_t l_pid; /* Process holding the lock. */
};
```

```
struct sembuf
{
    unsigned short int sem_num; /* semaphore number */
    short int sem_op; /* semaphore operation */
    short int sem_flg; /* operation flag */
};
```

```
typedef union
{
    char __size[__SIZEOF_SEM_T];
    long int __align;
} sem_t;
```

DEMO!

© M. Rashid Zamani



Fork

```
std::stringstream DATA;

void iWantToExecuteThisInTheChild() {
    uint8_t i = 0;
    while(++i<3) {
        std::cout << "-----" << std::endl;
        std::cout << ">>> pid = " << getpid() << std::endl;
        DATA << (getpid()) << std::endl;
        std::cout << DATA.str() << std::endl;
        std::cout << "-----" << std::endl << std::endl;
        sleep(1);
    }
}
```

```
std::cout << "My process id = " << getpid() << std::endl;
signal(SIGCHLD, SIG_IGN);

pid_t pid;
pid = fork();
if (pid);
else {
    for (size_t i=0 ; i < 2 ; i++) {
        pid = fork();

        /* if ( pid ) {
            break;
        }*/
        std::cout << "Child #" << getpid() << std::endl;
    }
    iWantToExecuteThisInTheChild();
}
```

Signal

```
if (pid == 0) { /* child */
    signal(SIGHUP, sighup);
    signal(SIGINT, sigint);
    signal(SIGQUIT, sigquit);
    for (;;)
        ; /* loop for ever */
}
```

```
else /* parent */
{ /* pid hold id of child */
    printf("\nPARENT: sending SIGHUP\n\n");
    kill(pid, SIGHUP);

    sleep(3); /* pause for 3 secs */
    printf("\nPARENT: sending SIGINT\n\n");
    kill(pid, SIGINT);

    sleep(3); /* pause for 3 secs */
    printf("\nPARENT: sending SIGQUIT\n\n");
    kill(pid, SIGQUIT);
    sleep(3);
}
```

Pipe

```
int pipeFileDescriptos[2];  
char buffer[30];  
  
pipe(pipeFileDescriptos);
```

```
if (!fork()) {  
    printf(" CHILD: writing to the pipe\n");  
    for (size_t i = 0; i < 5; i++) {  
        //read(pipeFileDescriptos[0], buffer, 5);  
        write(pipeFileDescriptos[1], "child", 5);  
        read(pipeFileDescriptos[0], buffer, 5);  
        printf("CHILD: read \"%s\"\n", buffer);  
    }  
  
    printf(" CHILD: exiting\n");  
    exit(0);  
} else {  
    printf("PARENT: reading from pipe\n");  
    for (size_t i = 0; i < 5; i++) {  
        write(pipeFileDescriptos[1], "test", 5);  
        read(pipeFileDescriptos[0], buffer, 5);  
        printf("CHILD: read \"%s\"\n", buffer);  
    }  
    printf("PARENT: read \"%s\"\n", buffer);  
    wait(NULL);  
}
```

FIFO

```
char s[200];  
int ret, fd;  
  
mknod(FifoName, S_IFIFO | 0644, 0);
```

```
fd = open(FifoName, O_RDONLY | O_WRONLY);  
ret = read(fd, s, 200);  
ret = write(fd, s, strlen(s));
```